

Scaling Language Models

Antoine Bosselut

Announcements


- **Guest Lecture Tomorrow!**
- **Proposal Feedback ready by end of week!**
- **Course Project:** Milestone 1 due this Sunday!
 - Data Packages were released. If you haven't received them yet, contact us ASAP

Today's Outline

- **Lecture**

- **Quick Recap:** Scale
- **Managing scale when training:** Scaling laws
- **Managing scale when deploying:** Model Compression, Speculative Decoding
 - how can we make LLMs more compute- and memory-efficient for deployment ?

Language Model Scaling



Larger models

More data

More compute

More




Every part of the model scales!

| Model Name | n_{params} | n_{layers} | d_{model} | n_{heads} | d_{head} | Batch Size | Learning Rate |
|-----------------------|---------------------|---------------------|--------------------|--------------------|-------------------|------------|----------------------|
| GPT-3 Small | 125M | 12 | 768 | 12 | 64 | 0.5M | 6.0×10^{-4} |
| GPT-3 Medium | 350M | 24 | 1024 | 16 | 64 | 0.5M | 3.0×10^{-4} |
| GPT-3 Large | 760M | 24 | 1536 | 16 | 96 | 0.5M | 2.5×10^{-4} |
| GPT-3 XL | 1.3B | 24 | 2048 | 24 | 128 | 1M | 2.0×10^{-4} |
| GPT-3 2.7B | 2.7B | 32 | 2560 | 32 | 80 | 1M | 1.6×10^{-4} |
| GPT-3 6.7B | 6.7B | 32 | 4096 | 32 | 128 | 2M | 1.2×10^{-4} |
| GPT-3 13B | 13.0B | 40 | 5140 | 40 | 128 | 2M | 1.0×10^{-4} |
| GPT-3 175B or “GPT-3” | 175.0B | 96 | 12288 | 96 | 128 | 3.2M | 0.6×10^{-4} |

- Trained on 570 GB of Common Crawl data
- **How?** Used cluster provided by Microsoft

Why scale?




- Last week, we talked about benefits of scaling in terms of **emergence**
- Practically, training for longer also leads to lower test loss
- Larger models can reach lower test losses

What should we scale?

Model size, dataset size, compute budget

Given a compute budget, how big of a model can we train?
and how big of a dataset should we train it on?

Impact of compute budget




**Dotted lines estimate these curves.
Need to predict for larger models!**

- For a fixed compute budget, there is an optimal number of parameters that we can train
- Having **too large** a model for **too small** a compute budget does NOT let the model learn
 - Model doesn't see enough examples during training
- Having **too small** a model for too large a compute budget is also bad
 - Repeated computation isn't helpful if the model has no capacity to encode additional information

Consideration


- With a fixed compute budget (in FLOP-days), we have two costs:
 - Number of floating point operations needed to train on a single example (model size)
 - Number of total examples we will train on
- **How should we trade off these two costs?**
 - Which should we prioritise?

Model-Data Trade-off




- Larger models benefit more from larger datasets
- **Smaller models saturate**
 - Only so much capacity to learn!
- At some point, larger models don't benefit more from same-sized data
- Model size needs to be scaled jointly with data size

Other Cool Findings




No need to make models terribly deep




Multiple ratios of depth vs. width (aka embedding size) are fine

Other cool findings



- LSTMs also follow scaling laws, benefitting from increased scale
- They scale less efficiently than transformers, though
- They still have trouble modelling long-term dependencies (>100 tokens)

To scale up: estimate model, data, compute




- Assuming no bottlenecks, expected test loss has power law relationship with each variable
- From smaller models, we can estimate how much compute, data, and model size is needed to achieve a particular test loss

Model Scaling in the last few years

| Model | Size (# Parameters) | Training Tokens |
|--|---------------------|-----------------|
| LaMDA (Thoppilan et al., 2022) | 137 Billion | 168 Billion |
| GPT-3 (Brown et al., 2020) | 175 Billion | 300 Billion |
| Jurassic (Lieber et al., 2021) | 178 Billion | 300 Billion |
| <i>Gopher</i> (Rae et al., 2021) | 280 Billion | 300 Billion |
| MT-NLG 530B (Smith et al., 2022) | 530 Billion | 270 Billion |

What happens if we get
these estimates wrong?

Oops!



- Chinchilla authors found that original works on model scaling had poorly estimated power laws
- New estimates showed that a 4x smaller model should be used for the compute budget
- Trained Gopher (280B) before finding this out!

Model Scaling in the last few years

| Model | Size (# Parameters) | Training Tokens |
|--|---------------------|-----------------|
| LaMDA (Thoppilan et al., 2022) | 137 Billion | 168 Billion |
| GPT-3 (Brown et al., 2020) | 175 Billion | 300 Billion |
| Jurassic (Lieber et al., 2021) | 178 Billion | 300 Billion |
| <i>Gopher</i> (Rae et al., 2021) | 280 Billion | 300 Billion |
| MT-NLG 530B (Smith et al., 2022) | 530 Billion | 270 Billion |
| <i>Chinchilla</i> | 70 Billion | 1.4 Trillion |

**Chinchilla gets better performance than all of the above models
on most common NLP benchmarks!**

Smaller model, trained on much more data!

Should we train the largest model that will
converge given the data and compute we have ?
(i.e., following Chinchilla scaling laws)

Not necessarily ! Why not ?

Inference cost!

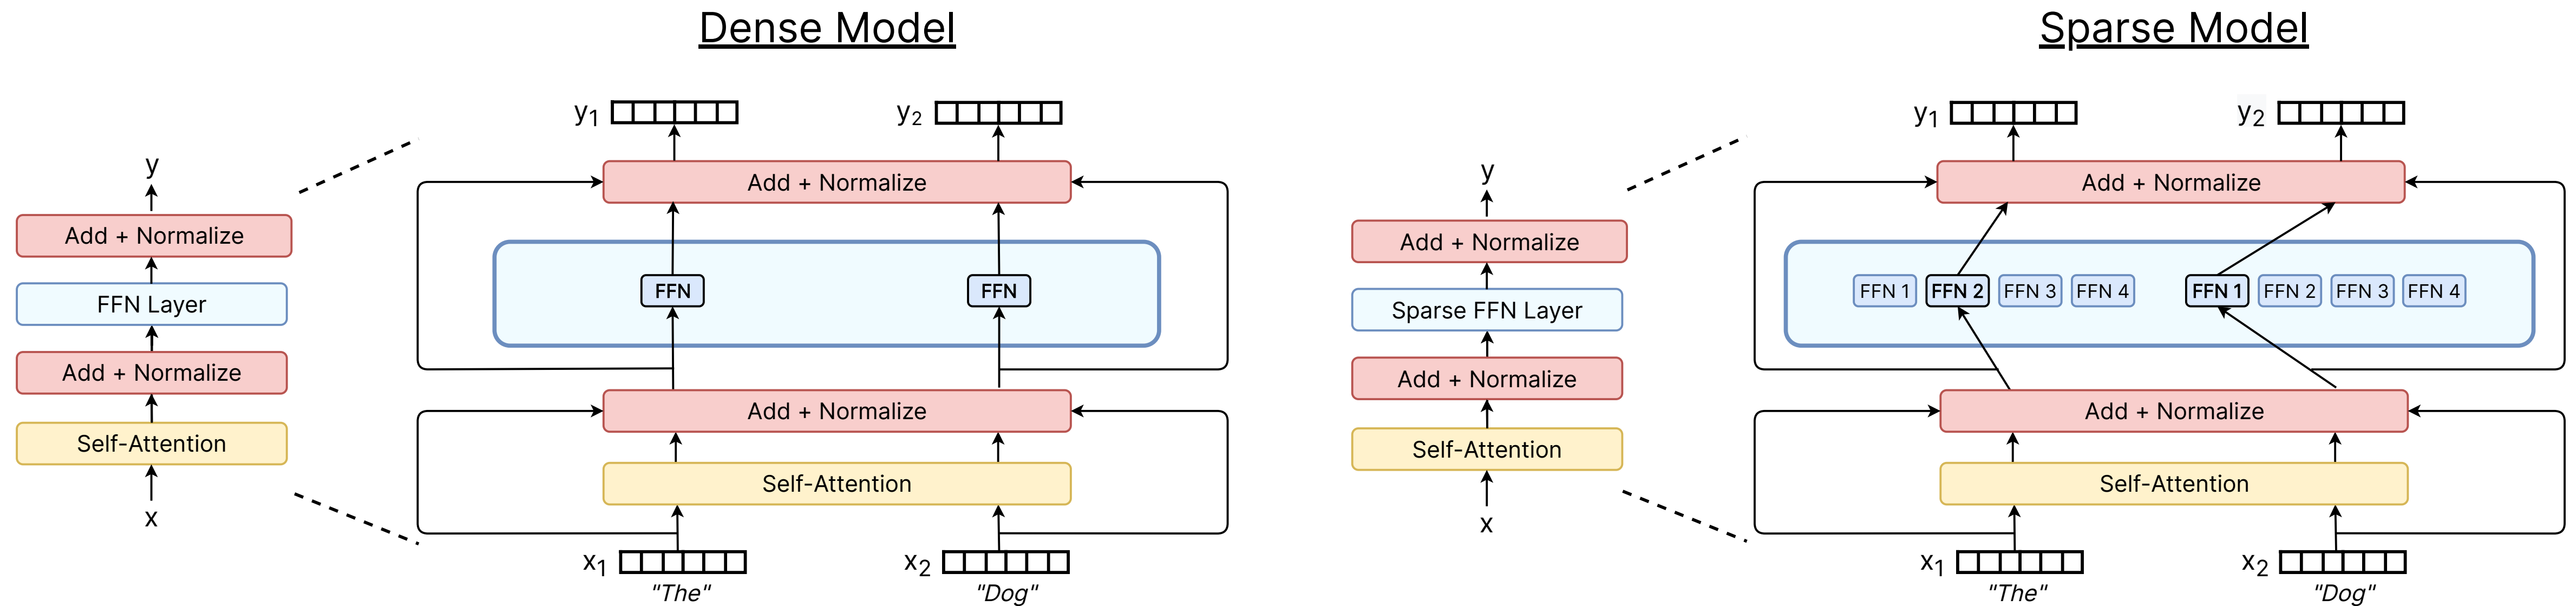
Importance of Inference

| | GPU Type | GPU Power consumption | GPU-hours | Total power consumption |
|------------|-----------|-----------------------|-----------|-------------------------|
| OPT-175B | A100-80GB | 400W | 809,472 | 356 MWh |
| BLOOM-175B | A100-80GB | 400W | 1,082,880 | 475 MWh |
| LLaMA-7B | A100-80GB | 400W | 82,432 | 36 MWh |
| LLaMA-13B | A100-80GB | 400W | 135,168 | 59 MWh |
| LLaMA-33B | A100-80GB | 400W | 530,432 | 233 MWh |
| LLaMA-65B | A100-80GB | 400W | 1,022,362 | 449 MWh |

- Scaling laws helps estimate dataset and model size for a given *training compute budget*
 - Ignores, the compute *inference budget*
 - How much should a single query cost ?
 - Training cost is amortised; inference cost is constant
- LLaMa authors showed that training smaller models (7B) on more data (1T tokens) continued to improve them
- Worse performance than 65B model, but much cheaper for inference (10x!)

How can we reduce inference cost while still keeping model capacity high ?

Mixture-of-Experts



- Initialise multiple **FFNs** in the transformer block
- Initialise routing function that selects an **FFN** that the out of self-attention should be routed to
 - Input can be routed to multiple **FFNs** (i.e., Top-K routing), but top-2 is common
- Model can have more parameters as number of "experts" increases, but inference cost per example remains the same

GPT-4, DeepSeek are mixture-of-experts architectures

Recap

- Scale is necessary to achieve many of the emergent breakthroughs of the last few years
 - in-context learning, chain-of-thought reasoning, instruction learning
- Training at scale is very expensive
 - Potentially, months of training = millions of \$\$\$\$
- Scaling laws let us estimate the optimal model and dataset sizes for a fixed compute budget, so that we only have to do the training once!
- While scaling laws suggests we should train the largest model possible, downstream *inference cost* is important to consider as well
 - Next module: **Compression!**

References

- Kaplan, J., McCandlish, S., Henighan, T., Brown, T. B., Chess, B., Child, R., ... & Amodei, D. (2020). Scaling laws for neural language models. arXiv preprint arXiv:2001.08361.
- Hoffmann, J., Borgeaud, S., Mensch, A., Buchatskaya, E., Cai, T., Rutherford, E., ... & Sifre, L. (2022). Training compute-optimal large language models. arXiv preprint arXiv:2203.15556.